

GROWTH AUDIT

Growth, Risk & IPO Readiness

A consolidated management pack on sustainable activation, automated financial-crime defence and investigation priority



EXECUTIVE SIGNALS

71.7% Sustainable activation proxy vs 78.2% reconstructed headline	GIB Highest-probability corridor 15.68% confirmed-fraud rate
84.65% Fraud after KYC pass 12,310 of 14,543 events	5 Priority investigations 17.71% of all fraud events

What this management pack answers

Three linked questions, one consistent transaction dataset and one control narrative.

ASSIGNMENT

Assess whether growth is sustainable, identify where financial-crime risk concentrates, and translate the results into actions that support safe scale, regulatory credibility and IPO readiness.

01 · Growth quality

Recalculate the claimed app conversion rate using a stricter definition tied to repeat economic behaviour. Explain the discrepancy and propose an investor-grade KPI architecture.

02 · Risk strategy

Separate geographic attack volume from fraud probability. Analyse fraud after KYC pass and design continuous, explainable defences with Swiss FINMA considerations.

03 · Investigation priority

Identify the five highest-priority fraud-linked users without relying only on event count or recorded amount. Challenge the cutoff and define the reserve queue.

Executive outputs

71.7% Sustainable proxy repeat card use; no confirmed fraud	GIB Highest probability GBR remains highest volume	84.65% Post-KYC fraud KYC is not durable trust	5 + 1 Priority queue five cases; one immediate reserve
--	---	---	---

How senior readers should use the pack

Reader	Primary decision	Relevant pages
Executive Committee / Board	Approve metric definitions, risk appetite, control ownership and the 90-day plan.	3–6, 10, 14
Growth / Product / Finance	Replace a one-event headline with cohort-based economic activation and separate risk quality.	3–6
Risk / Compliance / Operations	Deploy dual geographic lenses, continuous behavioural controls and tiered investigations.	7–14

EVIDENCE BASELINE

All quantitative findings trace to the supplied transaction file: 688,651 rows and 8,021 users. The file has no timestamps, realised-loss field, device or network links, or final investigation outcomes; recommendations are therefore management analysis, not automatic customer decisions.

01

A Defensible Conversion Rate

Recalculation, regulatory alignment and IPO-readiness implications

DECISION REQUESTED

Use 71.7% as the stricter current-file proxy for sustainable behaviour in this audit, while stating clearly that it is not a true app-funnel rate. Retire the unqualified “78% App Conversion Rate” from external or investor materials until cohort, window and event logic are controlled and reproducible.

78.2%

Marketing rate reconstructed
5,463 / 6,989 · ≥1 card-payment row


75.3%

Matured one-event quality
exclude users with confirmed fraud

71.7%

Stricter repeat-use proxy
≥2 card rows; no confirmed fraud

Marketing's 78% is reproducible; risk maturation lowers it



Bottom line

Marketing’s figure is arithmetically reproducible, but the label is too broad. The 78.2% calculation selects 6,989 users whose observed KYC value is always PASSED and counts any user with one CARD_PAYMENT row. The stricter result removes 200 users with confirmed fraud and 249 additional users without repeat card use — 449 users, or 6.42 percentage points, on the same denominator. The stricter result is closer to sustainable behaviour because it requires recurrence and a basic risk-quality screen. The dataset still lacks app installs, sign-up starts, timestamps, settlement status, reversals, refunds, chargebacks and a fixed fraud-maturity window. Every one of the 8,021 users already has a transaction. Neither 78.2% nor 71.7% is therefore a true app-conversion, retention or profitability metric.

MANAGEMENT POSITION

Keep 71.7% as an internal audit sensitivity. Build a prospective 90-day cohort KPI for repeat, settled economic use, and report financial-crime quality as a separate matured outcome.

The discrepancy is a definition choice—not a calculation error

The denominator stays fixed; the quality gates change.

Definition	Users	Rate	Δ pp	What changes
Marketing reconstruction	5,463	78.17%	—	One observed card payment.
Matured one-event quality	5,263	75.30%	−2.86	Remove 200 users with any confirmed fraud.
Stricter repeat-use proxy	5,014	71.74%	−6.42	Then remove 249 users with fewer than two card payments.
Strict proxy / all file users	5,014	62.51%	−15.66	Change denominator to all 8,021 observed users.

Marketing uses a late-stage denominator that excludes 1,032 users whose observed KYC status is FAILED, NONE or PENDING, then accepts one card event as “conversion.” The stricter rate holds the 6,989-person denominator constant; the 6.42-point decline comes only from confirmed-fraud screening and repeat observed use.

Acceptable wording “In the supplied transaction extract, 71.7% of users observed only as KYC PASSED had at least two card-payment rows and no confirmed fraud.”	Wording to avoid “71.7% of Revolut app users sustainably converted.” The required app-funnel denominator and event history are absent.
---	--

Research evidence: from first use to retained economic quality

Evidence	Finding	Implication for Revolut
Gupta & Zeithaml	Retention and lifetime-value metrics connect customer behaviour to financial performance.	Track downstream persistence and value; do not stop at one card event.
Gupta, Lehmann & Stuart	Customer value depends on future earnings, retention, margin and acquisition.	Pair activation with risk-adjusted contribution CLV, acquisition cost and payback.
Reinartz & Kumar	Long-life customers are not necessarily profitable.	Repeat use is necessary evidence, but not sufficient proof, of value.
Nu Holdings comparator	Monthly activity is defined by revenue generation in the prior 30 days.	Use precise, time-bounded, economically anchored definitions.

EBA remote-onboarding guidance requires risk-sensitive due diligence, pre-implementation assessment, ongoing monitoring and fraud-triggered review. FCA Consumer Duty material similarly requires firms to understand outcomes across the customer journey. A conversion target must not reward removal of necessary KYC, fraud or comprehension controls.

KEY DESIGN DISTINCTION Later confirmed fraud is essential quality evidence, but it should not silently rewrite the historical behavioural activation rate. Lock the cohort and report a separate lagged financial-crime quality result using the same cohort.

A two-metric architecture for sustainable growth

Separate behavioural activation from the matured financial-crime outcome.

1 · SEAR-90 — primary behavioural KPI For monthly cohort m, eligible members must show at least two positive-value, customer-initiated, settled qualifying transactions on distinct local days during days 0–30; at least one external card purchase or payment; and one further qualifying settled transaction during days 61–90. Publish $SEAR-90(m) = A(m) / C(m)$ only after the window and latency buffer close.	2 · MFCQR-90 — separate lagged guardrail $F(m)$ is the number of SEAR-90 achievers with no pre-specified, objectively adjudicated and timestamped financial-crime or account-abuse outcome through the day-90 lock date. $MFCQR-90(m) = F(m) / A(m)$. Report confirmed cases, pending cases and detection lag; never rewrite SEAR-90 silently.
GOVERNANCE REQUIREMENT The outcome taxonomy requires sign-off from Financial Crime, Compliance, Legal, Conduct, Data Protection and Internal Audit.	

Minimum data contract and control set

Control	Required evidence
Cohort	Legal entity, country, product, KYC approval, eligibility and person-level de-duplication timestamps.
Transaction lineage	Authorisation, settlement, reversal, refund, chargeback, value, currency, local date and a closed taxonomy.
Risk ledger	Decision-time onboarding snapshot plus later case-open, adjudication, confirmation and correction timestamps.
Calculation	Versioned deterministic code, frozen inputs, reconciliation, exception tests, reviewer and approval evidence.
Disclosure	Purpose, formula, counts, maturity, scope, limitations, comparatives, methodology changes and restatement policy.
Outcome safeguards	Complaints, accessibility, friction and lawful group-outcome monitoring shown beside—not netted into—the KPI.

Disclosure logic

Behavioural statement “A defined cohort completed repeat, settled economic use within the stated window.”	Risk-quality statement “Of those activated users, a defined share remained free of matured adjudicated financial-crime outcomes through the lock date.”
Prohibited shortcut Do not collapse both metrics into one historical percentage or imply that one card event establishes a primary banking relationship.	Restatement discipline Record methodology changes, detection lag and corrections. Preserve the originally published cohort result and explain any restatement.

Control the growth metric before using it in an equity story

A 90-day build converts an audit sensitivity into a governed, testable KPI.

VALUE THESIS

Revolut's FY2025 disclosures cite 68.3 million retail customers, £4.5 billion revenue, £1.7 billion profit before tax and £1.3 trillion transaction volume. A controlled activation KPI strengthens those disclosures by linking eligible cohort → repeat settled use → retention and contribution → visible risk and conduct guardrails.

90-day implementation plan

Timing	Management action	Decision gate
0–30 days	Freeze the 78% query, dashboard export, source population and wording. Approve SEAR-90/MFCQR-90 definitions and ownership.	Marketing and Legal withdraw or relabel the claim if documentary substantiation is incomplete.
31–60 days	Build immutable event lineage and the same-cohort bridge. Reconcile counts; test missingness, duplicates, ordering and settlement logic.	Finance Controls and Risk explain material bridge items; Internal Audit agrees the assurance approach.
61–90 days	Run shadow cohorts by country, channel and product. Test whether SEAR-90 predicts day-90/12-month activity, contribution, losses and outcomes.	The disclosure committee locks the definition only after predictive and conduct review; otherwise keep it internal.

IPO-readiness value chain

Comparable cohorts A stable definition across periods makes trends credible and enables controlled explanation of methodology changes.	Risk-adjusted growth Separating activation from matured financial-crime quality prevents promotional growth from obscuring residual risk.
Assurance evidence Lineage, reconciliations, reviewers, approvals and restatement rules make the KPI independently reperformable.	Disclosure restraint Do not publish SEAR-90 or MFCQR-90 as a headline IPO KPI until multiple comparable cohorts and assurance evidence exist.

BOARD RECOMMENDATION

Approve 71.7% as the auditable answer to this exercise with the stated limitations. Approve SEAR-90 and MFCQR-90 for prospective build and validation. Do not present either as a headline external KPI until the evidence base is mature.

02

Automating Defences for Growth

Geographic exposure, post-KYC behaviour and a FINMA-aligned control plan

EXECUTIVE VERDICT

Gibraltar is the highest-probability corridor; Great Britain is the largest fraud-volume pool. Run two lenses at once: rate-sensitive corridor controls for probability outliers and capacity/loss controls for high-volume markets. KYC must be a starting risk state—not a durable trust label.

15.68%

GIB fraud rate
138 / 880 · 95% CI 13.43–18.23%

6,631

GBR fraud volume
45.60% of all fraud events

84.65%


Fraud after KYC pass
12,310 of 14,543 events

260


Passed-KYC users
require behavioural review

Part A — one country drives volume; another drives probability

High volume: where attacks accumulate



High probability: where each transaction is riskier



High volume · GBR

6,631 fraud-labelled events from 146,200 transactions: 45.60% of the fraud-event pool, but a 4.54% fraud probability. Use for capacity, queue design, loss budgeting and concentration reporting.

High probability · GIB

138 fraud-labelled events from 880 transactions: 15.68%, or 7.43× the 2.11% global rate and 3.46× GBR. Use for corridor thresholds, step-up authentication, holds and focused investigation.

SINGLE HIGHEST-RISK COUNTRY: GIBRALTAR

This conclusion uses confirmed-fraud probability among merchant-country corridors with at least 50 transactions. The 95% Wilson interval is 13.43%–18.23%, and GIB remains the leader with a minimum exposure threshold of 500.


Data-quality warning: MERCHANT_COUNTRY = UNKNOWN contains 205,578 rows and 5,464 fraud events. It is not a country; it is a material lineage and control-completeness issue. Overall, 29.85% of rows have unknown merchant country.

Passing identity checks is no longer a durable trust signal

Fraud after KYC pass shares a repeat cash-movement and transfer signature.

12,310 Passed-KYC fraud events 84.65% of all fraud events	260 Distinct users KYC PASSED + fraud activity	23.5 Median recurrence fraud events per selected user
--	---	--

Post-KYC fraud is cash-cycle-heavy, round-numbered and amount-extreme



Shared behavioural signature

- 1.86× cash-cycle share
- 4.14× bank-transfer share
- 1.70× amounts divisible by 100
- 1.87× amounts divisible by 1,000
- 2.63× top-decile amount within type and currency

Persistence and breadth

- 76.15% use at least three methods
- 80.77% span at least two merchant countries
- 100.5 events at the 90th percentile
- 26.13% of events in the top 10 users
- 89.42% of fraud rows are GBP-denominated

INTERPRETATION

The cohort behaves less like ordinary card-led spending and more like repeat cash movement and transfer activity. Roundness and large relative amounts add structure signals. Geography and currency are exposure features—not proxies for nationality.

Control warning: “round to 1,000” is a subset of “round to 100”; production features must be encoded hierarchically to avoid double counting. The file has no timestamps, so velocity and sequence remain untested.

Replace static trust with continuous, explainable defence

Start with transparent rules; instrument decisions and outcomes before increasing model complexity.

1. **Corridor probability control.** Trigger on merchant-country fraud rate above an approved tolerance, with minimum n and confidence-bound guardrails. Raise risk, request step-up authentication or hold for review; do not auto-decline on geography alone.

2. **Post-KYC behavioural score.** Combine cash-cycle share, amount percentile, hierarchical roundness, recurrence and method/country breadth. Use review tiers, reason codes, customer-friction limits and investigator override.

3. **Data-quality fallback.** When merchant country, amount or type/currency mapping is missing, route to enrichment, apply a conservative fallback and prevent silent score degradation.

4. **Human and model control.** On material change, drift, backlog breach, override spike or adverse customer outcome, pause or roll back, document root cause and retest before release.

Target operating model

1 · Ingest	2 · Score	3 · Decide	4 · Learn
Transaction · KYC · device · beneficiary	Corridor · behaviour · network	Allow · step-up · hold · review	Outcome · loss · override · drift

Technology path

Approach	Trade-offs	Cost	Set-up
Rules-first layer	Fast and explainable; weaker against adaptation; best immediate control.	Low–medium	4–8 weeks
Hybrid rules + model	Better interaction detection; requires labels, validation, drift and override controls.	Medium	8–16 weeks
Network/device layer	Strong for linked abuse; requires new data, privacy design and integration.	Medium–high	12–24 weeks

RECOMMENDED ROUTE

Launch a rules-first minimum-viable control now, instrument every decision and outcome, then add supervised and network models behind approval gates.

Control metrics before automation scales

Detection Confirmed fraud, realised loss, missed-event findings and precision/recall where outcomes exist.	Operations Queue age, backlog, review SLA, investigator capacity and manual-override rate.
Customer outcomes Step-up completion, complaints, friction, accessibility and lawful segment outcomes.	Resilience Data completeness, control availability, rollback, recovery objectives and severe-scenario tests.

FINMA-aligned governance turns automation into a scalable control

Circulars define supervisory expectations; they do not prescribe a named fraud engine or threshold.

SCOPE GUARDRAIL Applicability depends on the Swiss licence, FINMA category, booking model and exemptions. Confirm the final control mapping with Swiss counsel and the applicable official-language text.		
Authority	Requirement	Design evidence
FINMA Circ. 2023/1 § 23–38	Board/executive risk governance, risk inventory, tolerance, indicators and independent testing.	Named owners, appetite, KRIs/KCIs, model/rule/vendor inventory and independent test results.
FINMA Circ. 2023/1 § 49–112	Controlled ICT change, critical-data controls, BCM, resilience and severe-scenario testing.	Release approvals, lineage, access reviews, backups, manual fallback and recovery objectives.
FINMA Circ. 2017/1 § 6–97	Board oversight, management information, ICS and independent risk/compliance/audit challenge.	Three-lines accountability, board reporting, issue ownership, ageing and closure evidence.
FINMA Circ. 2016/7 § 5–53	Video/online identity proofing, evidence, escalation and third-party controls at onboarding.	Keep as an onboarding control—not a continuous behavioural-monitoring mandate.
AMLO-FINMA Art. 20	Effective monitoring of business relationships and transactions.	Risk-based monitoring, investigation and documented escalation/reporting decisions.

90-day control plan

When	Deliverable	Board-ready evidence
0–30 days	Rules-first pilot	GIB corridor, UNKNOWN-geography fallback, behavioural reasons, human override and baseline false-positive sample.
31–60 days	Control governance	Risk appetite, ownership, inventory, change gate, independent test plan, manual fallback and incident playbook.
61–90 days	Scale decision	Out-of-time test, loss/friction/capacity metrics, segment calibration, remediation log and IPO-quality workpapers.

Value created

Safer scale Risk-adjusted launch gates allow products and markets to grow without silently accumulating loss, backlog or compliance exposure.	Operational credibility Evidence of ownership, validation, change control and resilience makes the defence repeatable and auditable.
Disclosure credibility Reconcile prevented fraud, realised loss, residual exposure, customer friction and control failures.	Investor confidence The equity story becomes controlled, repeatable risk-adjusted growth—not a promotional prevented-fraud headline.

03

The Five Accounts to Investigate First

A behaviour-led priority queue for scalable defence and IPO readiness

EXECUTIVE DECISION

Open coordinated investigations on five masked accounts now. Ninety per cent of the score measures repeatability, statistical conviction, transaction-mix abnormality and attack breadth. Type/currency-normalised amount severity is capped at ten per cent.

299

Eligible fraud-linked users
at least one fraud-labelled event

2,575

Top-5 fraud events
17.71% of all fraud events


5 / 5

Full method coverage
every selected target

3 / 5

KYC passed
identity did not end risk

Five evidence-rich investigation priorities



Why these five

Every selected target combines a 100% observed fraud label, high event persistence and all five transaction methods. The group spans 3–11 merchant countries, separating campaign capability from a one-off large transaction.

Decision boundary

The fifth and sixth scores differ by only 1.12 points. Keep rank 6 in an immediate reserve queue; the score supports triage, not accusation or automatic restriction.

Method and weights

Repeatability	Conviction	Abnormality	Breadth	Severity
30% · log event count	25% · Wilson lower bound	20% · excess vs mix	15% · methods/markets/currencies	10% · top-three amount percentiles

Calculation chain: 1 Eligibility → 2 Normalise within type × currency → 3 Combine five explicit weighted signals → 4 Challenge with counterfactuals and sensitivity. Demographics do not enter the score; mixed-currency totals are excluded.

Why each account made the cut

Masked identifiers remain unique within the supplied file.


#1 · dc283b17...71d9 · 99.4 1,029 events · five methods · 11 markets · severity rank 1 · KYC PENDING. Dominates repeatability, conviction, abnormality and severity. Preserve evidence and map linked devices, beneficiaries and funding sources first.	#2 · b8271606...a49f · 97.8 340 events · five methods · three markets · severity rank 6 · KYC PASSED. High conviction with four currencies. Test coordinated cross-currency movement and multiple funding channels.
#3 · 25c2ecb3...1217 · 97.6 460 events · five methods · five markets · severity rank 31 · KYC PASSED. Strongest excess fraud versus expected transaction mix; moderate amount severity proves risk is not reducible to value.	#4 · 5c75c857...e8de · 96.6 238 events · five methods · three markets · severity rank 20 · KYC PENDING. Balanced profile with high severity and full-method reach; verify the extreme-event evidence explicitly.
#5 · 4ee8690a...edc6 · 96.0 508 events · five methods · eight markets · severity rank 84 · KYC PASSED. Low severity rank is outweighed by recurrence, conviction, abnormality and broad five-method activity.	

Case-opening protocol

FIRST-LINE ACTION Open coordinated fraud, anti-money-laundering and cyber investigations. Preserve transaction and onboarding evidence; enrich with devices, IPs, beneficiaries, counterparties, funding sources, prior alerts and linked-account graphs.	
Evidence preservation Freeze relevant transaction, onboarding, authentication, device and case records. Record scope, custodian, time and access.	Link analysis Test shared devices, IPs, beneficiaries, counterparties, funding sources and merchant-country overlaps.
Customer treatment Apply legal, proportionality and customer-outcome review before restriction, closure or reporting decisions.	Decision record Keep model version, reason codes, reviewer, override, evidence and final disposition for independent re-performance.
BOUNDARY These are investigation-priority targets—not findings of identity, criminal conduct or legal liability.	

Why higher event count or amount severity does not settle the queue

The cutoff is defensible only when its trade-offs and near-ties are explicit.



Nearest excluded · 6eb51e3f...c930

590 events and 18 markets exceed selected #5, but the pattern covers four methods rather than five. Breadth scores 85.90 versus 95.71; composite scores are 94.89 versus 96.01. Decision: immediate reserve and concurrent review if capacity permits.

Amount challenger · 1632f5e0...63d3

Amount-severity score 99.67 and severity rank 2, but only 41 events, three methods and one market. Composite rank is 91 with a score of 68.05. Decision: retain in a separate high-severity queue.

Sensitivity: the first three are robust; the boundary is policy-sensitive

Target	Base	No amt	Amt 25%	Equal	No width	Conv.
dc283b17...71d9	1	1	1	1	1	1
b8271606...a49f	2	5	2	2	2	3
25c2ecb3...1217	3	3	3	3	3	2
5c75c857...e8de	4	7	4	4	4	5
4ee8690a...edc6	5	2	7	5	7	4

GOVERNANCE RULE

Treat score differences below three points as approximate ties. Preserve the component table, model version and reason codes so an investigator can reperform every selection.

Immediate operating decision: open the five selected cases, preserve evidence and place rank 6 into shadow review. Maintain the amount challenger in a separate high-severity queue rather than allowing one large value to displace a broader campaign target.

Make the queue auditable, resilient and board-ready

Swiss supervisory expectations, a 90-day evidence plan and transparent AI-assisted authorship.

Authority	Requirement	Queue evidence
FINMA Circ. 2023/1 § 23–38	Risk governance, inventory, tolerance, indicators and independent testing.	Board appetite, owner, model/rule inventory, weight/cutoff rationale, validation and remediation log.
FINMA Circ. 2023/1 § 49–82	Controlled ICT change, critical-data confidentiality, integrity, availability and access.	Source-to-score lineage, approved releases, access reviews, pseudonymised reports, audit logs and rollback.
FINMA Circ. 2023/1 § 83–112	Business continuity and resilience testing.	Manual fallback, recovery objectives, dependency map and fraud-surge/bad-data tests.
FINMA Circ. 2017/1 § 6–97	Board oversight, ICS and independent risk/compliance/internal-audit challenge.	Three-lines ownership, documented challenge, issue ageing, closure evidence and board reporting.
AMLO-FINMA Art. 20	Effective ongoing monitoring.	Risk-based monitoring, investigation and documented escalation/reporting decisions.

90-day evidence plan

When	Deliverable	Board-ready evidence
0–10 days	Investigation pack	Five cases plus reserve #6; evidence holds, reason codes, linkage map and legal/customer-treatment review.
11–45 days	Model governance	Approved weights and cutoff, challenger tests, independent review, overrides, issue log and rollback.
46–90 days	Scale and disclosure pack	Time/device/network/loss enrichment; yield, false positives, friction, realised loss and residual-risk reconciliation.

Authorship and methodology disclosure

AI-ASSISTED REPORT

This document was produced with AI assistance. Analysis, calculations, chart generation and drafting were carried out from the supplied dataset and consolidated at the user's request. All figures should be independently reperformed and reviewed by a qualified analyst before external, regulatory or board disclosure.

The calculations, definitions and figures follow a documented, versioned methodology. The repository contains the exact queries, cohort definitions, scoring logic, dashboard and reproducible analysis used across the three tasks.

Methodology repository: github.com/Razije/revolut-fincrim-growth-risk-audit

Primary regulatory references

[FINMA Circular 2023/1](#) · [FINMA Circular 2017/1](#) · [FINMA Circular 2016/7](#) · [AMLO-FINMA](#)

IMPORTANT NOTICE

This management analysis is not legal, tax, accounting or investment advice. Percentages are rounded; unrounded values were used for differences. The work does not constitute Swiss legal advice, an IPO prediction, a venue recommendation, a valuation conclusion or a timetable.